

# Random forest taxonomic classifier of Insecta with k-mer feature selection from CytB sequences

Nathaniel Lesperance

2023-11-05

```
library(rentrez)
library(tidyverse)
library(Biostrings)
library(randomForest)
library(DECIPHER)
library(RSQLite)
library(ggplot2)
library(caret)
library(pROC)
```

## Set Variables and Parameters

```
# Set variables to be used throughout script
TAXON1 <- "Coleoptera" # 1st Taxonomy of interest
TAXON2 <- "Lepidoptera" # 2nd Taxonomy of interest

# Max proportion of Ns allowed in sequences
prop_Ns_cutoff <- 0.02

# Gene of interest (name, approx length)
GENE_OI <- c("CytB", "800:1200")

# Random forest downstream parameters
num_trees <- 50
kmer <- 8
```

## Code Section I:

Create a helper function to be used for efetching FASTA sequences

```
# Query NCBI with taxon and gene of interest,
# using esearch and efetch to get FASTA files, collecting only DNA sequences and
# write them to a fasta file and read them into a DNASTringSet to be used downstream
efetch_FASTAs <- function(geneOI, taxon) {
  # Preliminary nucleotide db esearch
```

```

goi_search <- entrez_search(db = "nuccore",
                           term = paste(taxon, "[ORGN] AND",
                                         geneOI[1], "[Gene] AND",
                                         geneOI[2], "[SLEN] AND biomol_genomic[PROP]"))

# Repeat esearch with retmax as count from esearch result above
# and use history = T for downstream efetch
goi_search <- entrez_search(db = "nuccore",
                           term = paste(taxon, "[ORGN] AND",
                                         geneOI[1], "[Gene] AND",
                                         geneOI[2], "[SLEN] AND biomol_genomic[PROP]"),
                           retmax = goi_search$count,
                           use_history = T)

Sys.sleep(5) # Wait 5s so don't query NCBI too frequently

# Get FASTA sequences from NCBI using efetch,
# using web history from esearch result above
goi_fastas <- entrez_fetch(db = "nuccore",
                          rettype = "FASTA",
                          web_history = goi_search$web_history)

# Report results
print(paste(goi_search$count,
            "DNA FASTA sequences fetched matching",
            geneOI[1],
            "from Taxonomic Group",
            taxon))

# Write fetched sequences to fasta file then read into DNASTringSet
write(goi_fastas, paste(taxon, "_", geneOI[1], ".fasta"), sep = "\n")
taxon1_seq <- readDNASTringSet(paste(taxon, "_", geneOI[1], ".fasta"))

return(taxon1_seq)
}

```

## Import Sequences from NCBI

```

# eFetch FASTA sequences that match taxonomic groups
tax1_seq <- efetch_FASTAs(geneOI = GENE_OI, taxon = TAXON1)

## [1] "1139 DNA FASTA sequences fetched matching CytB from Taxonomic Group Coleoptera"

tax2_seq <- efetch_FASTAs(geneOI = GENE_OI, taxon = TAXON2)

## [1] "3376 DNA FASTA sequences fetched matching CytB from Taxonomic Group Lepidoptera"

# Preliminary examination of sequences
BrowseSeqs(tax1_seq)
BrowseSeqs(tax2_seq)
# Some shorter sequences

```

```

# Name sequences and build df's from them with taxon label
tax1_df <- data.frame(Taxonomy = TAXON1,
                      Title = names(tax1_seq),
                      Sequence = paste(tax1_seq))
tax2_df <- data.frame(Taxonomy = TAXON2,
                      Title = names(tax2_seq),
                      Sequence = paste(tax2_seq))

# Parse species name from FASTA label
tax1_df$Species <- word(tax1_df$Title, start = 2L, end = 3L)
tax2_df$Species <- word(tax2_df$Title, start = 2L, end = 3L)

# Build quick table to check descriptors for sequences
table(word(tax1_df$Title, start = -1, end = -1))

```

```

##
##      Contig5 mitochondrial      product
##      1          1090          48

```

```
table(word(tax2_df$Title, start = -1, end = -1))
```

```

##
##      cds mitochondrial
##      11          3365

```

```

# Report number of unique species
sprintf("%d initial unique species of %s", length(unique(tax1_df$Species)), TAXON1)

```

```
## [1] "236 initial unique species of Coleoptera"
```

```
sprintf("%d initial unique species of %s", length(unique(tax2_df$Species)), TAXON2)
```

```
## [1] "160 initial unique species of Lepidoptera"
```

## Create a helper function to be used for Quality Control (QC)

```

# Function to filter gaps, Ns and by length
quality_and_length_filt <- function(tax_df, taxon) {

  # Report and save initial length distribution stats
  summ_stats <- summary(str_count(tax_df$Sequence))
  print(paste(taxon, "length distribution before QC"))
  print(summ_stats)

  print(paste(sum(str_count(tax_df$Sequence, "-")),
              "gaps and", sum(str_count(tax_df$Sequence, "N")),
              "Ns in all", taxon,
              "sequences were detected before QC"))
}

```

```

# Save quartiles for downstream length filtering
q1 <- summ_stats[2]
q3 <- summ_stats[5]
IQR <- q3 - q1

# Save number of sequences before QC
before_length <- nrow(tax_df)

tax_df <- tax_df %>%
  # Remove starting and ending gaps or Ns
  mutate(Seq2 = str_remove_all(Sequence, "^N+|N+$|-")) %>%
  # Remove sequences with greater than 2% Ns
  filter(str_count(Seq2, "N") <= (prop_Ns_cutoff * str_count(Sequence))) %>%
  # Remove sequences with length >1.5 times IQR above 3rd quartile or
  # length >1.5 times IQR below 1st quartile
  filter(str_count(Seq2) >= (q1 - 1.5 * IQR) &
         str_count(Seq2) <= (q3 + 1.5 * IQR))

# Report length distribution stats after QC
print(paste(taxon, "length distribution after QC"))
print(summary(str_count(tax_df$Sequence)))

print(paste(sum(str_count(tax_df$Seq2, "-")),
            "gaps and", sum(str_count(tax_df$Seq2, "N")),
            "Ns in all", taxon,
            "sequences were detected after QC"))

# Report number of sequences before and after QC
print(paste(taxon, "sequences before QC:", before_length))
print(paste(taxon, "sequences after QC:", nrow(tax_df)))

return(tax_df)
}

```

```

# QC on both sets of sequences
tax1_df_clean <- quality_and_length_filt(tax1_df, TAXON1)

```

```

## [1] "Coleoptera length distribution before QC"
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      807   1049   1134   1057   1137   1161
## [1] "0 gaps and 88 Ns in all Coleoptera sequences were detected before QC"
## [1] "Coleoptera length distribution after QC"
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      922   1122   1137   1126   1137   1161
## [1] "0 gaps and 15 Ns in all Coleoptera sequences were detected after QC"
## [1] "Coleoptera sequences before QC: 1139"
## [1] "Coleoptera sequences after QC: 886"

```

```

tax2_df_clean <- quality_and_length_filt(tax2_df, TAXON2)

```

```

## [1] "Lepidoptera length distribution before QC"

```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      825.0  961.0   991.0   977.4   994.0  1152.0
## [1] "0 gaps and 31 Ns in all Lepidoptera sequences were detected before QC"
## [1] "Lepidoptera length distribution after QC"
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      913.0  977.0   991.0   982.8   994.0  1035.0
## [1] "0 gaps and 31 Ns in all Lepidoptera sequences were detected after QC"
## [1] "Lepidoptera sequences before QC: 3376"
## [1] "Lepidoptera sequences after QC: 2959"
```

```
# Create dfs for QC plotting purposes
len1df_B <- tax1_df %>%
  mutate(Source = paste(TAXON1, "Before")) %>%
  mutate(Seq2 = Sequence)

len2df_B <- tax2_df %>%
  mutate(Source = paste(TAXON2, "Before")) %>%
  mutate(Seq2 = Sequence)

len1df_A <- tax1_df_clean %>%
  mutate(Source = paste(TAXON1, "After"))

len2df_A <- tax2_df_clean %>%
  mutate(Source = paste(TAXON2, "After"))

length_df <- rbind(len1df_B, len2df_B, len1df_A, len2df_A)

# Plot before and after sequence length distributions
ggplot(data = length_df, mapping = aes(x = as.vector(str_count(Seq2)), fill = Source)) +
  geom_density(alpha = 0.4, n = 512) +
  labs(x = "Sequence Length (bp)",
       y = "Density",
       title = sprintf("%s and %s Sequence Length Distribution Before and After QC",
                       TAXON1, TAXON2)) +
  theme_bw()
```

## Coleoptera and Lepidoptera Sequence Length Distribution Before and Aft

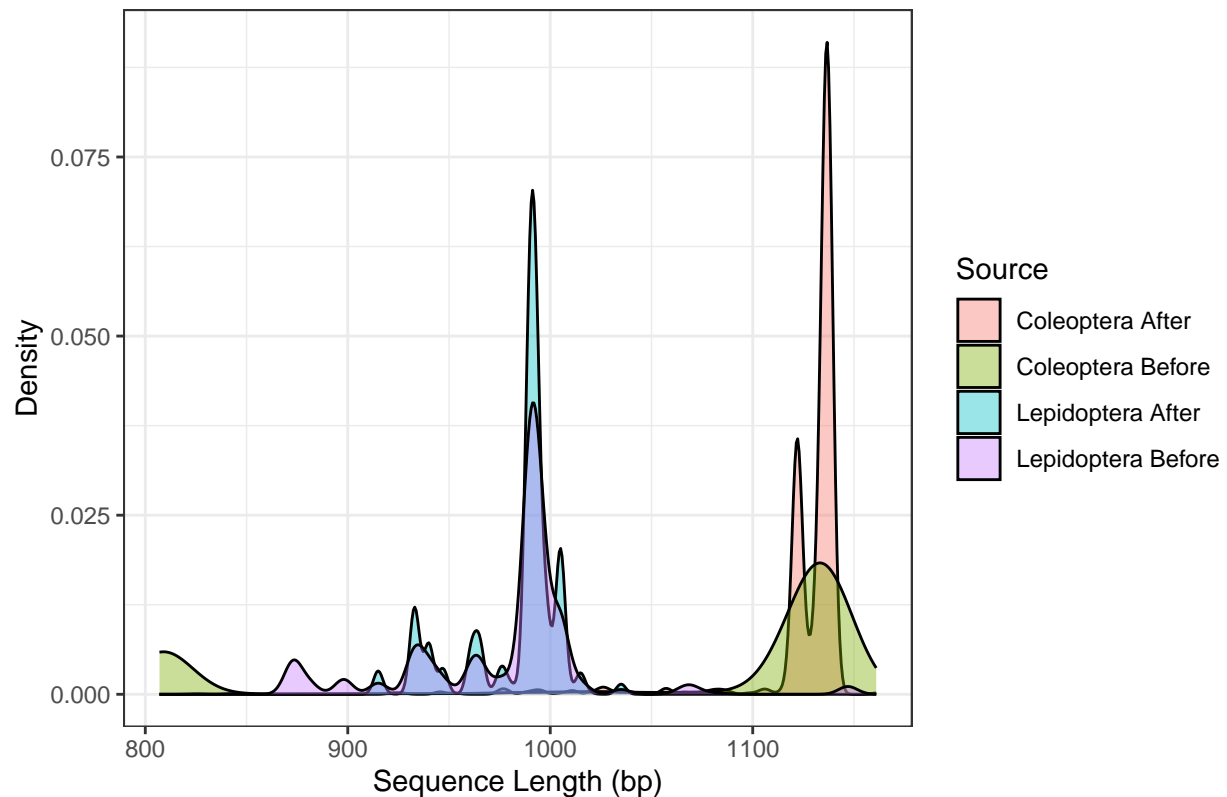


Figure 1: Sequence length distributions for each taxon before and after quality control (QC) filtering. The Coleoptera inter-quartile range of sequence length before QC was 88 bp and after it was 15 bp. The Lepidoptera inter-quartile range of sequence length before QC was 33 bp and after it was 17 bp.

```
# Examine Sequences after QC
BrowseSeqs(DNAStringSet(tax1_df_clean$Seq2))
BrowseSeqs(DNAStringSet(tax2_df_clean$Seq2))

# Ensure using same number of sequences from each
min_seq <- min(dim(tax1_df_clean)[1], dim(tax2_df_clean)[1])

set.seed(42)
tax1_df_clean <- tax1_df_clean %>%
  sample_n(min_seq)

set.seed(42)
tax2_df_clean <- tax2_df_clean %>%
  sample_n(min_seq)

# Report number of unique species
sprintf("%d unique species of %s after QC", length(unique(tax1_df$Species)), TAXON1)

## [1] "236 unique species of Coleoptera after QC"
```

```
sprintf("%d unique species of %s after QC", length(unique(tax2_df$Species)), TAXON2)
```

```
## [1] "160 unique species of Lepidoptera after QC"
```

## Code Section II:

### Extract k-mer Features from Sequences

```
all_taxa <- rbind(tax1_df_clean, tax2_df_clean)

features_df <- data.frame(Taxonomy = all_taxa$Taxonomy,
                          Name = all_taxa$Title,
                          oligonucleotideFrequency(DNAStringSet(all_taxa$Seq2),
                                                    as.prob = T, width = kmer),
                          verbose = TRUE)

# Check results of previous steps
table(features_df$Taxonomy)
```

```
##
## Coleoptera Lepidoptera
##      886      886
```

### Build Initial RF Classifier

```
# Create Validation set
set.seed(42)
valid_df <- features_df %>%
  group_by(Taxonomy) %>%
  sample_n(floor(0.2 * nrow(features_df)/2))

# Check equal representation
table(valid_df$Taxonomy)
```

```
##
## Coleoptera Lepidoptera
##      177      177
```

```
# Create Training set
set.seed(42)
train_df <- features_df %>%
  filter(!Name %in% valid_df$Name)
```

```
set.seed(42)
ranfor_classifier <- randomForest(x = train_df[, 3:ncol(train_df)],
                                y = as.factor(train_df$Taxonomy),
```

```

ntree = num_trees,
importance = T)

# Check that each row has been left out of bag sufficient number of times
# (To check if effective number of trees chosen)
summary(ranfor_classifier$oob.times)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      8.00   16.00   18.00   18.31   20.00   32.00

# Check initial classifier performance
ranfor_classifier$confusion

##              Coleoptera Lepidoptera class.error
## Coleoptera           709           0           0
## Lepidoptera           0           709           0

# Get predictions from training and validation set to get more detailed performance stats
predict_val_initial <- predict(ranfor_classifier,
                              valid_df[, 3:ncol(train_df)])
predict_train_initial <- predict(ranfor_classifier,
                                train_df[, 3:ncol(train_df)])

# Get detailed performance stats to report
confusionMatrix(predict_train_initial,
                 as.factor(train_df$Taxonomy))

## Confusion Matrix and Statistics
##
##              Reference
## Prediction  Coleoptera Lepidoptera
## Coleoptera           709           0
## Lepidoptera           0           709
##
##              Accuracy : 1
##              95% CI : (0.9974, 1)
##      No Information Rate : 0.5
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 1
##
## Mcnemar's Test P-Value : NA
##
##              Sensitivity : 1.0
##              Specificity : 1.0
##      Pos Pred Value : 1.0
##      Neg Pred Value : 1.0
##              Prevalence : 0.5
##      Detection Rate : 0.5
##      Detection Prevalence : 0.5
##      Balanced Accuracy : 1.0
##

```



```
##      'Positive' Class : Coleoptera
##
```

```
confusionMatrix(predict_val_initial,
                  as.factor(valid_df$Taxonomy))
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction   Coleoptera Lepidoptera
##  Coleoptera      177         0
##  Lepidoptera      0         177
##
##              Accuracy : 1
##              95% CI : (0.9896, 1)
##      No Information Rate : 0.5
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 1
##
##  Mcnemar's Test P-Value : NA
##
##      Sensitivity : 1.0
##      Specificity : 1.0
##      Pos Pred Value : 1.0
##      Neg Pred Value : 1.0
##      Prevalence : 0.5
##      Detection Rate : 0.5
##      Detection Prevalence : 0.5
##      Balanced Accuracy : 1.0
##
##      'Positive' Class : Coleoptera
##
```

```
# Use Mean Decrease in Accuracy to get measure of importance for each feature
feat_importance <- importance(ranfor_classifier, type = 1)

# Create df of features and their importance to classification
feat_importance <- data.frame(Feature = rownames(feat_importance),
                              Importance = feat_importance)

# Order by descending importance
feat_importance <- feat_importance[order(-feat_importance$MeanDecreaseAccuracy),]
```

```
# Examine 30 most important
head(feat_importance, n = 30)

# Examine 30 least important
tail(feat_importance, n = 30)

# Create character vector of features with importance above 0
refined_features <- feat_importance$Feature[feat_importance$MeanDecreaseAccuracy > 0]
length(refined_features) # 411 features
```

## Train new models using fewer features

```
# Set seed for consistent randomization
set.seed(42)

# Create vectors for varying number of features included in the models
num_features <- c(seq(1, 100, 1))
err_rate1 <- vector("numeric", length(num_features))
err_rate2 <- vector("numeric", length(num_features))
val_rate <- vector("numeric", length(num_features))

# Loop through the different number of features, gradually adding features
for (i in 1:length(num_features)) {
  set.seed(42)
  ranfor_classifier_refined <- randomForest(x = train_df[, refined_features[1:(i+1)]],
                                            y = as.factor(train_df$Taxonomy),
                                            ntree = num_trees,
                                            importance = F)

  # Save error rate of last iteration of RF algorithm
  err_rate1[i] <- ranfor_classifier_refined$err.rate[num_trees,2]
  err_rate2[i] <- ranfor_classifier_refined$err.rate[num_trees,3]

  # Predict classification on validation set
  predict_val <- predict(ranfor_classifier_refined,
                        valid_df[, refined_features[1:(i+1)]])

  # Save validation error rate
  val_rate[i] <- mean(predict_val == valid_df$Taxonomy)
}

# Create df for downstream graphing error rate over number of features included in model
graph_err_df <- data.frame(Features = num_features,
                          error_taxon1 = err_rate1,
                          error_taxon2 = err_rate2,
                          val_err_rate = (1-val_rate))

# Plot error rate from identification of each taxon during training
# along with combined error in validation
ggplot(graph_err_df, aes(x = Features)) +
  geom_smooth(aes(y = error_taxon1, color = paste(TAXON1, "Training"))) +
  geom_smooth(aes(y = error_taxon2, color = paste(TAXON2, "Training"))) +
  geom_smooth(aes(y = val_err_rate, color = "Validation (Mean)")) +
  scale_color_manual(values = c("orange", "black", "blue")) +
  labs(
    title = "Error Rates vs. Features Included in RF Model",
    x = "Number of Features",
    y = "Class Error Rate",
    color = "Dataset Type") +
  theme_bw()

## 'geom_smooth()' using method = 'loess' and formula = 'y ~ x'
## 'geom_smooth()' using method = 'loess' and formula = 'y ~ x'
## 'geom_smooth()' using method = 'loess' and formula = 'y ~ x'
```




Figure 2: Class error rate from Training on each taxonomy of interest and mean class error rate from Validation, over number of features included in the model. Shading indicates 95% confidence interval of class error.

```
# Create plot showing 20 most important features (motifs)
ggplot(feats_importance[order(-feats_importance$MeanDecreaseAccuracy)[1:20],],
  aes(x = MeanDecreaseAccuracy, y = Feature)) +
  geom_point(size = 3, color = "blue") +
  labs(title = "20 Most Impactful DNA Motifs for Classification",
    x = "Mean Decrease in Model Accuracy", y = "Motifs") +
  theme_bw()
```

20 Most Impactful DNA Motifs for Classification

